We have scaling laws in terms of compute and spatial resolution of our data (represented by Voxel size). We plot loss on the y-axis and voxel size / number of steps on the x axis.

We see log-linear scaling with spatial resolution at small spatial resolutions, and no returns for spatial resolution at very large resolutions. However, the scale of current noninvasive imaging methods is within the log linear domain, giving good evidence that finetuning neural foundation models would allow them to generalize.

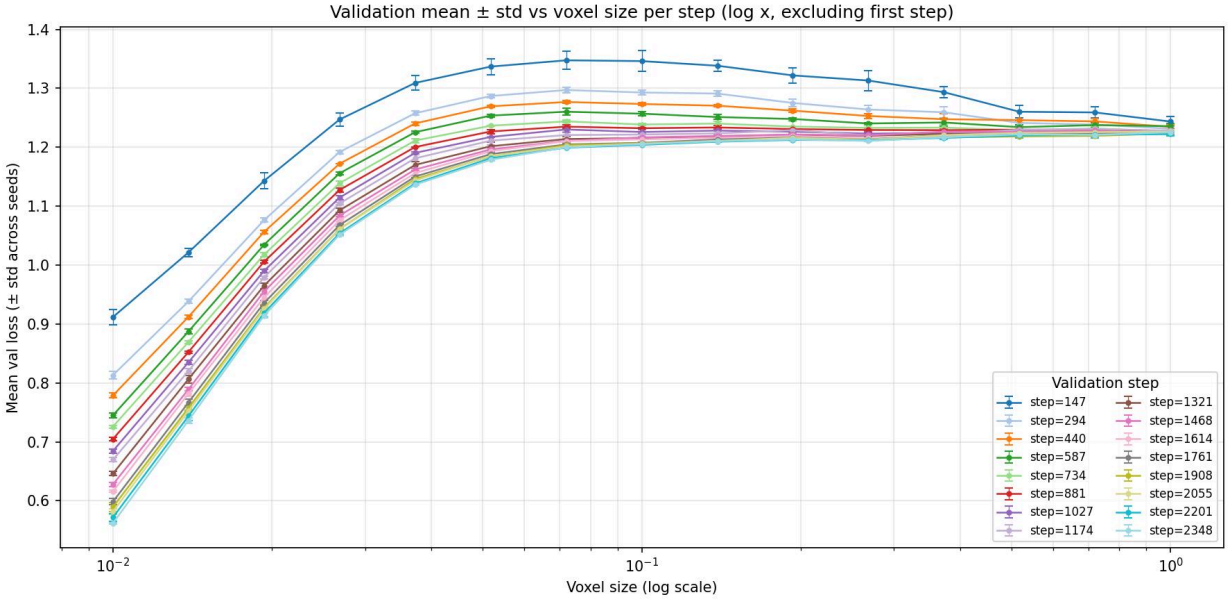

In addition, we see that we are slightly undertraining our models. At very low voxel sizes, where we have perfect neuron by neuron information, we are still solidly in a log-linear regime, which shows that we will still be able to get significant dividends to scaling.