

Sudarsh Kunnnavakkam

+1 (949) 254-8232 | Pasadena, CA | kvsudarsh786@gmail.com | github.com/skunnnavakkam | sudarsh.com

WORK EXPERIENCE

CTO, Co-Founder

Nov 2025 — Present

Cognitogram Labs

San Francisco, CA

- Co-founded venture-backed neuroscience startup pursuing whole-brain emulation; raised \$1M pre-seed from Character Capital at 18.
- Built an imaging facility from the ground up to image the whole brain of the zebrafish at 10,000 slices per second in 3 months
- Designed and trained autoregressive models of larval zebrafish neural activity toward whole-brain emulation.

Research Assistant

Sep 2023 — Jun 2025

Model Evaluation and Threat Research (METR)

Berkeley, CA

- Lead engineer for internal project to estimate the agentic time horizon of LLMs at much lower cost, through distilling failure points in agentic transcripts
- Co-lead engineer of a evaluation for Chain-of-Thought Faithfulness of Large Language Models, finding that Chain of Thought is informative about LLM cognition as long as the cognition is complex enough that it can't be performed in a single forward pass
- Helped lead teams of contractors red-team LLMs and curate datasets such as [DAFT Math](#) of difficult, free-response questions

Undergraduate Research Intern

Nov 2024 — Aug 2025

ShapiroLab at Caltech

Pasadena, CA

- Built a high-throughput ultrasound screening platform; designed custom proteins with RFDiffusion, AlphaFold, and ESM3 for faster ultrasound reporter kinetics.

Research Fellow

Feb 2025 — May 2025

Supervised Program for Alignment Research

Remote

- Implemented a *continuous double auction* agent arena as a model environment for LLM collusion, accepted to *ICML 2025*

EDUCATION

California Institute of Technology

Pasadena, CA

B.S. in Physics & Computer Science

In progress, on leave

- Co-director of Caltech AI Alignment, lead reading groups, research, and invited speakers; scaled from 10 to 100+ attendees at events

WRITING & OTHER WORK

- [Simulated Qualia Mugging](#)
- [Takes on Automating Alignment](#)
- [How Should We Implement an AI Pause?](#)
- [Redundant Attention Heads in Large Language Models For In Context Learning](#)

SELECTED PUBLICATIONS

- A. Deng*, S. Von Arx*, B. Snodin, [S. Kunnnavakkam](#), T. Lanham, "CoT May Be Highly Informative Despite "Unfaithfulness"" by *METR*, Aug 2025
- K. Agarwal, V. Teo, J. Vazquez, [S. Kunnnavakkam](#), V. Srikanth, A. Liu, "Evaluating LLM Agent Collusion in Double Auctions" at *ICML 2025 Workshop on Multi-Agent Systems in the Era of Foundation Models*, Vancouver, Canada, July 2025.
- C. J. Effarah*, T. Chen*, [S. Kunnnavakkam](#)*, C. M. Gonzalez, H. W. Lee, "Liquid Metal Printed 2D ITO for Nanophotonic Applications," in *California-US Government Workshop on 2D Materials*, Irvine, California, USA, Sep 2023

PROJECTS

METR: Faithfulness and Monitorability Eval 2025
• Co-lead engineer on METR research report on chain-of-thought (CoT) faithfulness (Aug 2025), extending Anthropic’s seminal evaluation to three frontier models and publishing findings for the wider safety community

LLM Agent Collusion Arena 2025
• Implemented a continuous double auction system for agents
• Implemented oversight, monitors, and other experimental conditions to test influence on collusion
• Added logging and metrics with WandB
• Accepted to ICML 2025 Workshop on Multi-agent Systems

Scanning Tunneling Microscope 2024
• Built working STM for \$1,000 using open-source design
• Achieved atomic-resolution imaging capabilities (Circuit Design, Signal Processing, Mechanical Engineering)

AWARDS

ARENA 6.0 Attendee 2025
Non-trivial Fellow 2024
YC W26 Acceptance (Declined) 2026